

1. Create a table named `Employees` with the following columns:

- `employee_id` (primary key, auto-increment)
- `first_name` (not null)
- `last_name` (not null)
- `email` (unique, not null)
- `hire_date` (default to the current date)
- `salary` (greater than or equal to 3000)

Ans: CREATE TABLE `Employees` (

`employee_id` INT PRIMARY KEY AUTO_INCREMENT,

`first_name` VARCHAR(50) NOT NULL,

`last_name` VARCHAR(50) NOT NULL,

`email` VARCHAR(100) UNIQUE NOT NULL,

`hire_date` DATE DEFAULT CURRENT_DATE,

`salary` DECIMAL(10, 2) CHECK (salary >= 3000)

);

2. Create a sequence `employee_seq` that starts at 1001 and increments by 1. Use this sequence to insert a new employee record into the `Employees` table.

-- Create the sequence

CREATE SEQUENCE `employee_seq`

START WITH 1001

INCREMENT BY 1;

-- Insert a new employee using the sequence

INSERT INTO `Employees` (`employee_id`, `first_name`, `last_name`, `email`, `salary`)

VALUES (`employee_seq.NEXTVAL`, 'John', 'Doe', 'john.doe@example.com', 5000);

3. Create a view `EmployeeView` that shows the `employee_id`, `first_name`, `last_name`, and `salary` of employees who have a salary greater than or equal to 5000.

Ans:

```
CREATE VIEW EmployeeView AS

SELECT employee_id, first_name, last_name, salary

FROM Employees

WHERE salary >= 5000;
```

4. Create an index on the `email` column of the `Employees` table to speed up search queries based on email addresses.

```
CREATE INDEX idx_email ON Employees (email);
```

5. Create a synonym `Emp` for the `Employees` table to simplify referencing the table in future queries.

Ans: `CREATE SYNONYM Emp FOR Employees;`

6. Add a new table `Departments` with the columns:

- `department_id` (primary key)
- `department_name` (not null) Then, alter the `Employees` table to add a foreign key constraint that references `department_id` from the `Departments` table.

Ans:

-- Create Departments table

```
CREATE TABLE Departments (

    department_id INT PRIMARY KEY,

    department_name VARCHAR(100) NOT NULL

);
```

-- Alter Employees table to add foreign key constraint

ALTER TABLE Employees

ADD department_id INT,

ADD CONSTRAINT fk_department FOREIGN KEY (department_id) REFERENCES
Departments(department_id);

7. Alter the Employees table to add a new column phone_number of type VARCHAR(15).

Ans: ALTER TABLE Employees

ADD phone_number VARCHAR(15);

8. Drop the phone_number column from the Employees table.

Ans:

ALTER TABLE Employees

DROP COLUMN phone_number;

9. Drop the Departments table from the database.

Ans: DROP TABLE Departments;

10. Create a trigger update_salary that automatically updates the salary column in the Employees table to 6000 whenever a new employee is inserted with a salary less than 3000.

Ans:

CREATE OR REPLACE TRIGGER update_salary

BEFORE INSERT ON Employees

FOR EACH ROW

BEGIN

IF :NEW.salary < 3000 THEN

:NEW.salary := 6000;

END IF;

END;

11. Alter the `Employees` table to add a unique constraint on the `phone_number` column.

Ans: ALTER TABLE Employees

ADD CONSTRAINT unique_phone_number UNIQUE (phone_number);

12. Alter the `salary` column in the `Employees` table to increase its size to `DECIMAL(15, 2)`.

Ans:

ALTER TABLE Employees

MODIFY salary DECIMAL(15, 2);

13. Rename the `Employees` table to `Staff`.

Ans: rename employees to staff.

14. Create a view called `DepartmentSalarySummary` that shows the `department_id`, `department_name`, and the total salary of all employees in that department.

Ans:

CREATE VIEW DepartmentSalarySummary AS

SELECT d.department_id, d.department_name, SUM(e.salary) AS total_salary

FROM Employees e

JOIN Departments d ON e.department_id = d.department_id

GROUP BY d.department_id, d.department_name;

15. Create a composite index `idx_name_salary` on the `first_name` and `salary` columns in the `Employees` table to speed up queries that filter on both fields.

Ans:

CREATE INDEX idx_name_salary ON Employees (first_name, salary);

16. Create a sequence `course_seq` that starts at 1000, increments by 1, and has a minimum value of 1000 and a maximum value of 9999.

Ans:

```
CREATE SEQUENCE course_seq  
  
START WITH 1000  
  
INCREMENT BY 1  
  
MINVALUE 1000  
  
MAXVALUE 9999  
  
CYCLE;
```

17. Create a table `Course_Enrollments` to track student course enrollments. The table should have:

- `enrollment_id` (primary key)
- `student_id` (foreign key referencing `Students.student_id`)
- `course_id` (foreign key referencing `Courses.course_id`)

Ans:

```
CREATE TABLE Course_Enrollments (  
  
    enrollment_id INT PRIMARY KEY,  
  
    student_id INT,  
  
    course_id INT,  
  
    enrollment_date DATE DEFAULT CURRENT_DATE,  
  
    FOREIGN KEY (student_id) REFERENCES Students(student_id),  
  
    FOREIGN KEY (course_id) REFERENCES Courses(course_id)  
  
);
```

18. Create a check constraint on the `hire_date` column of the `Employees` table to ensure that no employee can have a hire date in the future.

Ans: ALTER TABLE Employees

```
ADD CONSTRAINT check_hire_date CHECK (hire_date <= CURRENT_DATE);
```

19. Drop the synonym Emp that was created for the Employees table.

Ans:

DROP SYNONYM Emp;

DML:

Create the following tables :

- **Students** (student_id, first_name, last_name, email, department_id, enrollment_date, phone_number)
- **Departments** (department_id, department_name)
- **Courses** (course_id, course_name, department_id, course_fee)
- **Enrollments** (enrollment_id, student_id, course_id, enrollment_date, grade)

1. Get a list of all students in the university, showing their first and last names, email, and department name.

Ans:

```
SELECT s.first_name, s.last_name, s.email, d.department_name
FROM Students s
JOIN Departments d ON s.department_id = d.department_id;
```

2. Insert a new student into the Students table.

Ans:

```
INSERT INTO Students (student_id, first_name, last_name, email, department_id,
enrollment_date, phone_number)
VALUES (1002, 'Alice', 'Smith', 'alice.smith@example.com', 1, CURRENT_DATE, '555-1234');
```

3. Update the email address of the student with student_id = 1002.

Ans:

```
UPDATE Students
SET email = 'alice.smith_new@example.com'
WHERE student_id = 1002;
```

4. Delete a student record from the Students table where student_id = 1002.

Ans:

```
DELETE FROM Students
WHERE student_id = 1002;
```

5. Retrieve all courses offered by the "Computer Science" department, including the course name and fee.

Ans:

```
SELECT c.course_name, c.course_fee
FROM Courses c
```

```
JOIN Departments d ON c.department_id = d.department_id
WHERE d.department_name = 'Computer Science';
```

6. Find the names of all students enrolled in the course "Database Systems."

Ans:

```
SELECT s.first_name, s.last_name
FROM Students s
JOIN Enrollments e ON s.student_id = e.student_id
JOIN Courses c ON e.course_id = c.course_id
WHERE c.course_name = 'Database Systems';
```

7. Insert a new enrollment for student `student_id = 1001` in the course `course_id = 201`, setting the enrollment date to today's date and giving the grade 'A'.

Ans:

```
INSERT INTO Enrollments (enrollment_id, student_id, course_id, enrollment_date, grade)
VALUES (10001, 1001, 201, CURRENT_DATE, 'A');
```

8. Update the grade of student `student_id = 1001` for the course `course_id = 201` to 'B+'.

Ans:

```
UPDATE Enrollments
SET grade = 'B+'
WHERE student_id = 1001 AND course_id = 201;
```

9. Get the number of students enrolled in each course along with the course name.

Ans:

```
SELECT c.course_name, COUNT(e.student_id) AS num_students
FROM Courses c
LEFT JOIN Enrollments e ON c.course_id = e.course_id
GROUP BY c.course_name;
```

10. Find all students who have not enrolled in any courses.

Ans:

```
SELECT s.first_name, s.last_name
FROM Students s
LEFT JOIN Enrollments e ON s.student_id = e.student_id
WHERE e.enrollment_id IS NULL;
```

11. Retrieve the `first_name`, `last_name`, and `email` of all students enrolled in the "Mathematics" department.

Ans:

```
SELECT first_name, last_name, email
FROM Students
```

```
WHERE department_id = (SELECT department_id FROM Departments WHERE department_name =
'Mathematics');
```

12. Find all courses where the `course_fee` is greater than 2000, displaying the course name and fee.

Ans:

```
SELECT course_name, course_fee  
FROM Courses  
WHERE course_fee > 2000;
```

13. Retrieve a list of all students who have not yet paid their fees. Assume there is a `fee_paid` column in the `Students` table (boolean: 0 for not paid, 1 for paid).

Ans:

```
SELECT first_name, last_name, email  
FROM Students  
WHERE fee_paid = 0;
```

14. Count the number of courses each student is enrolled in, and display the student's name and the number of courses they are enrolled in.

Ans:

```
SELECT s.first_name, s.last_name, COUNT(e.enrollment_id) AS num_courses  
FROM Students s  
LEFT JOIN Enrollments e ON s.student_id = e.student_id  
GROUP BY s.student_id;
```

15. Get the `first_name`, `last_name`, and `grade` of all students who are enrolled in the course "Advanced Databases."

Ans:

```
SELECT s.first_name, s.last_name, e.grade  
FROM Students s  
JOIN Enrollments e ON s.student_id = e.student_id  
JOIN Courses c ON e.course_id = c.course_id  
WHERE c.course_name = 'Advanced Databases';
```

16. Update the grade of all students enrolled in the course "Database Systems" to 'A+' if their current grade is 'B'.

Ans:

```
UPDATE Enrollments e  
SET grade = 'A+'  
WHERE e.course_id = (SELECT course_id FROM Courses WHERE course_name = 'Database Systems')  
AND e.grade = 'B';
```

17. Get a list of students who have completed the course "Data Structures" (i.e., those who have a grade other than NULL).

Ans:


```
SELECT s.first_name, s.last_name
FROM Students s
JOIN Enrollments e ON s.student_id = e.student_id
JOIN Courses c ON e.course_id = c.course_id
WHERE c.course_name = 'Data Structures' AND e.grade IS NOT NULL;
```

18. Find the department offering the course with the highest fee.

Ans:

```
SELECT d.department_name, c.course_name, c.course_fee
FROM Departments d
JOIN Courses c ON d.department_id = c.department_id
WHERE c.course_fee = (SELECT MAX(course_fee) FROM Courses);
```

19.

Insert multiple students into the `Students` table with their `student_id`, `first_name`, `last_name`, `email`, `department_id`, and `phone_number`.

Ans:

```
INSERT INTO Students (student_id, first_name, last_name, email, department_id, phone_number)
VALUES
(1003, 'Bob', 'Brown', 'bob.brown@example.com', 2, '555-5678'),
(1004, 'Charlie', 'Green', 'charlie.green@example.com', 3, '555-6789'),
(1005, 'Diana', 'White', 'diana.white@example.com', 1, '555-7890');
```

20. Delete students from the `Students` table who are not enrolled in any course (i.e., no entries in the `Enrollments` table).

Ans: DELETE FROM Students

```
WHERE student_id NOT IN (SELECT DISTINCT student_id FROM Enrollments);
```

MongoDB

1.Question:

Insert a document into the `products` collection with the following data:

- name: "Laptop"

- category: "Electronics"
- price: 1200
- stock: 50

Ans:

```
db.products.insertOne({
  name: "Laptop",
  category: "Electronics",
  price: 1200,
  stock: 50
});
```

2. Insert multiple documents into the `orders` collection. Each document should have the following structure:

- `orderId`: Unique order identifier
- `customerName`: The customer's name
- `items`: Array of products (each product has name and quantity)
- `orderDate`: Date of the order

```
db.orders.insertMany([
  {
    orderId: 101,
    customerName: "Alice",
    items: [{ name: "Laptop", quantity: 1 }, { name: "Mouse", quantity: 2 }],
    orderDate: new Date("2024-11-01")
  },
  {
    orderId: 102,
    customerName: "Bob",
    items: [{ name: "Phone", quantity: 1 }],
    orderDate: new Date("2024-11-02")
  }
]);
```

3. Query the `products` collection to find all documents where the `category` is "Electronics".

Ans: `db.products.find({ category: "Electronics" }).toArray();`

4. Query the `products` collection to find all products where the `price` is greater than 1000.

Ans:

`db.products.find({ price: { $gt: 1000 } }).toArray();`

5. Update the stock quantity of the product with `name "Laptop"` in the `products` collection to 45.

Ans:

```
db.products.updateOne(
  { name: "Laptop" }, // Filter condition
  { $set: { stock: 45 } } // Update operation
);
```

6. Update the stock of all products in the `Electronics` category by adding 10 to the existing stock.

Ans:

```
db.products.updateMany(
  { category: "Electronics" }, // Filter condition
  { $inc: { stock: 10 } } // Increment the stock by 10
);
```

7. Delete the order with `orderId 101` from the `orders` collection.

Ans:

`db.orders.deleteOne({ orderId: 101 });`

8. Remove all products from the `products` collection that belong to the `category "Electronics"`.

Ans: `db.products.deleteMany({ category: "Electronics" });`

9. Insert a document into the `users` collection with the following data:

- `name: "Sara Lee"`
- `email: "sara.lee@example.com"`
- `age: 28`
- `address: { city: "San Francisco", state: "CA" }`

Ans: `db.users.insertOne({`
 `name: "Sara Lee",`
 `email: "sara.lee@example.com",`
 `age: 28,`
 `address: { city: "San Francisco", state: "CA" }`
`});`

10. Insert multiple blog post documents into the `posts` collection. Each document should have:

- `title`: Post title
- `author`: Author's name
- `content`: Blog content
- `tags`: Array of tags
- `published`: Boolean indicating if the post is published

Ans:

```
db.posts.insertMany([
  {
    title: "Introduction to MongoDB",
    author: "John",
    content: "This is an introductory post about MongoDB.",
    tags: ["MongoDB", "Database", "NoSQL"],
    published: true
  },
  {
    title: "Advanced MongoDB Queries",
    author: "Jane",
    content: "This post discusses advanced MongoDB query techniques.",
    tags: ["MongoDB", "Queries", "Aggregation"],
    published: false
  }
]);
```

11. Query the `users` collection to find all users where `age` is greater than 30.

Ans:

```
db.users.find({ age: { $gt: 30 } }).toArray();
```

12. Query the `posts` collection to find all blog posts that are marked as `published: true`.

Ans:

```
db.posts.find({ published: true }).toArray();
```

13. Update the `address` of the user with `name: "Sara Lee"` to the following:

- `city: "Los Angeles"`
- `state: "CA"`

Ans:

```
db.users.updateOne(  
  { name: "Sara Lee" }, // Filter condition  
  { $set: { address: { city: "Los Angeles", state: "CA" } } } // Update operation  
);
```

14. Add a new tag `"Tutorial"` to the post with `title: "Introduction to MongoDB"` in the `posts` collection.

Ans:

```
db.posts.updateOne(  
  { title: "Introduction to MongoDB" }, // Filter condition  
  { $addToSet: { tags: "Tutorial" } } // Add the tag to the tags array  
);
```

Q. Write a DDL query to create a table called `users` with the following columns:

- `user_id` (Primary Key, Integer, Auto-increment)
- `name` (Varchar, 100 characters)
- `email` (Varchar, 100 characters, Unique)
- `age` (Integer)

Ans: CREATE TABLE `users` (

```
  user_id INT AUTO_INCREMENT PRIMARY KEY,
```

```
name VARCHAR(100),  
email VARCHAR(100) UNIQUE,  
age INT,  
);
```